Data Safari #3

Week 10 · Replicating a published figure

Same format as Data Safari #1 — find a paper, recreate a figure from its public data, and reflect on the process. Render this file to PDF and submit both the `.qmd` and the rendered `.pdf` via Canvas alongside this week's regular assignment, before Week 11's lecture (Monday). See the syllabus for the late policy.

Tip: with regression and ANCOVA under your belt, look for a figure showing a **relationship between two continuous variables**, ideally with multiple groups shown by color and a fitted trend line per group — the kind of figure an ANCOVA model would produce.

Instructions

1. **Find a paper** with a scatterplot-style figure and publicly available data.

Full citation:

2. **Recreate the figure**, including the fitted line(s), working from the actual data.
3. **Reflect**: how easy was the data to find and use? How closely does your figure match? What would you change about the paper's data documentation?

Bonus (optional)

If the paper reports a regression coefficient or R^2, try fitting the same model yourself and compare your numbers to theirs. Then use `avg_slopes()` (or `avg_comparisons()`) from the `marginaleffects` package to report the slope in each group — do your numbers match the paper's?